

whatpulse_viewer_v4_as_pdf

March 16, 2024

1 WhatPulse Keypress Stats Analyzer

By Kenneth Burchfiel

Released under the MIT license

(I am not affiliated with WhatPulse (<https://whatpulse.org>) but highly recommend trying out the program, which I've used since September 2008. You can find my online WhatPulse page here: <https://whatpulse.org/KBurchfiel>)

This script allows you to perform various analyses of your WhatPulse typing stats. It does so by accessing the local WhatPulse database on your computer (along with other copies of databases you might want to include); reading this data into Pandas DataFrames, and then summarizing and graphing that data. The output of the script includes:

1. Rolling average calculations at the 7-, 28-, and 365-day level
2. Percentile and rank calculations
3. Weekly and hourly keypress stats
4. Static (.png) and interactive (.html) keypress charts

By converting this notebook into a Python file and then instructing your computer to run it on an hourly basis, you can also keep track of how much you're typing during your day. Which is sort of nerdy, to be honest. But so is this whole program ;)

I'll first import a number of packages that the program will use:

```
[ ]: import time
start_time = time.time() # Allows the program's runtime to be measured
import pandas as pd
import sqlalchemy
import numpy as np
from datetime import date
import statsmodels.api as sm
from scipy.stats import percentileofscore
import plotly.express as px
import kaleido
from IPython.display import Image
import datetime
```

I'll next define the names of folders that will store various outputs, along with default settings to use when converting interactive charts to static .png files:

```
[ ]: static_graphs_folder = 'graphs/static/'
interactive_graphs_folder = 'graphs/interactive/'
data_folder = 'data'

default_image_height = 540
default_aspect_ratio = 16 / 9 # Standard HD/UHD aspect ratio
default_image_width = default_image_height * default_aspect_ratio
default_image_scale = 5 # Creating a smaller graph (e.g. one 540 pixels
# in height) and then scaling it helps keep the text a bit larger.
```

1.1 Importing Whatpulse data

In order to analyze my WhatPulse data, I'll first need to import it from my local Whatpulse SQLite database. I'll also import a copy of the Whatpulse SQLite database stored on my old laptop (so that my analysis doesn't have to be limited to my current computer's keypress data.)

You'll of course need to update the following cell with the path(s) to your own WhatPulse database(s).

Note: if you want to run this program on your own, but don't have access to a WhatPulse database, you can still run this program (assuming you've downloaded or cloned it from GitHub). Skip ahead to the line where I read in

```
[ ]: database_paths_list = ['C:/Users/kburc/AppData/Local/whatpulse/whatpulse.db',
'C:/Users/kburc/D1V1/Documents/whatpulse_database_backups/a13r2_whatpulse.db',
'G:/My Drive/whatpulse_database_backups/linux_whatpulse.db']
# Note that the first path is to my computer's active database,
# whereas additional paths point towards copies of databases from
# other computers that I have used.
```

The following function analyzes each database's table at either the daily or hourly level.

```
[ ]: def generate_keypress_totals(database_path, level = 'daily'):
    '''The level argument can be 'daily', in which case the DataFrame
    returned by the function will show daily keypress totals, or 'hourly',
    which will preserve the hourly keypress totals in the original database.'''
    file_name = database_path.split('/')[-1] # Retrieves the final element
    # of the path (e.g. the file name)
    sqlalchemy_sqlite_engine = sqlalchemy.create_engine(
        'sqlite:///'+database_path)
    # Based on https://docs.sqlalchemy.org/en/13/dialects/sqlite.
    ↪html#connect-strings
    sqlalchemy_connection = sqlalchemy_sqlite_engine.connect()
    df_keypresses = pd.read_sql("select * from keypresses",
    con = sqlalchemy_sqlite_engine) # Simply reads all of the data from this
    # Capitalizing column names so that less renaming will be necessary
    # when creating graphs:
    df_keypresses.columns = [column.title() for column in df_keypresses.columns]
```

```

# table into a Pandas DataFrame
df_keypresses = df_keypresses.query("Day != '0000-00-00']").copy() # Removes
# this blank date value from the database if it happens to be there
if level == 'daily': # In this case, we'll want to combine hourly keypress
    # totals into a single row for each day
    df_keypresses = df_keypresses.pivot_table(
        index = 'Day', values = 'Count', aggfunc = 'sum')
    df_keypresses.sort_values('Day', inplace = True)
elif level == 'hourly': # The original data is already displayed
    # at the hourly level, so there's no need for a pivot_table() call.
    df_keypresses.sort_values(['Day', 'Hour'], inplace = True)
else:
    raise ValueError("Unrecognized level argument passed to function.")
df_keypresses.rename(columns={'Count': 'Keypresses'}, inplace=True)
return df_keypresses

```

I'll now run `generate_keypress_totals` in order to create a record of daily keypresses for both my current database and a copy of a past database. (I'll look at hourly keypress totals later on.)

```

[ ]: sqlalchemy_sqlite_engine = sqlalchemy.create_engine(
    'sqlite:///'+ 'C:/Users/kburc/AppData/Local/whatpulse/whatpulse.db')
sqlalchemy_connection = sqlalchemy_sqlite_engine.connect()
sqlalchemy_connection

```

```

[ ]: <sqlalchemy.engine.base.Connection at 0x2c4e8ff77c0>

```

```

[ ]: keypress_databases_list = []

for path in database_paths_list: # This loop creates a DataFrame for
    # each WhatPulse database stored in database_paths_list.
    print("Now loading:", path)
    keypress_databases_list.append(generate_keypress_totals(
        path, level = 'daily'))

# I'll now combine these tables into a single DataFrame.
df_combined_daily_keypresses = pd.concat(
    [keypress_databases_list[i] for i in range(len(keypress_databases_list))])
df_combined_daily_keypresses.sort_index(inplace=True)

# At this point, my copy of df_combined_daily_keypresses has multiple
# entries for days in which I logged keys on multiple operating systems.
# Therefore, the following line groups these entries into a single row
# for each date.
df_combined_daily_keypresses = df_combined_daily_keypresses.reset_index(
).pivot_table(index = 'Day', values = 'Keypresses', aggfunc = 'sum')
df_combined_daily_keypresses.index = pd.to_datetime(
    df_combined_daily_keypresses.index)

```

```
df_combined_daily_keypresses.to_csv('data/df_combined_daily_keypresses.csv')
df_combined_daily_keypresses
```

Now loading: C:/Users/kburc/AppData/Local/whatpulse/whatpulse.db

Now loading:

C:/Users/kburc/D1V1/Documents/whatpulse_database_backups/a13r2_whatpulse.db

Now loading: G:/My Drive/whatpulse_database_backups/linux_whatpulse.db

[]: Keypresses

| Day | |
|------------|-------|
| 2020-09-21 | 19081 |
| 2020-09-22 | 32771 |
| 2020-09-23 | 32065 |
| 2020-09-24 | 34698 |
| 2020-09-25 | 47038 |
| ... | ... |
| 2024-03-12 | 65458 |
| 2024-03-13 | 57509 |
| 2024-03-14 | 67124 |
| 2024-03-15 | 36963 |
| 2024-03-16 | 29865 |

[1204 rows x 1 columns]

The following line rebuilds df_combined_daily_keypresses using a copy of the DataFrame that got exported to a .csv file earlier on. This cell allows allow you to run this script even if you don't have your own WhatPulse database.

```
[ ]: df_combined_daily_keypresses = pd.read_csv(
      'data/df_combined_daily_keypresses.csv', index_col='Day')
      # The following line makes the index compatible with
      # date operations that the following code block will perform.
      df_combined_daily_keypresses.index = pd.to_datetime(
      df_combined_daily_keypresses.index)
      df_combined_daily_keypresses
```

[]: Keypresses

| Day | |
|------------|-------|
| 2020-09-21 | 19081 |
| 2020-09-22 | 32771 |
| 2020-09-23 | 32065 |
| 2020-09-24 | 34698 |
| 2020-09-25 | 47038 |
| ... | ... |
| 2024-03-12 | 65458 |
| 2024-03-13 | 57509 |

| | |
|------------|-------|
| 2024-03-14 | 67124 |
| 2024-03-15 | 36963 |
| 2024-03-16 | 29865 |

[1204 rows x 1 columns]

The following code block fills in the DataFrame with missing dates (e.g. dates in which I did not have any keypresses). I want to add in those missing dates so that I can calculate more accurate rolling keypress averages.

```
[ ]: first_date = df_combined_daily_keypresses.index[0]
last_date = df_combined_daily_keypresses.index[-1]
full_date_range = pd.date_range(start=first_date, end = last_date)
# https://pandas.pydata.org/docs/reference/api/pandas.date\_range.html
df_combined_daily_keypresses = df_combined_daily_keypresses.reindex(
    full_date_range, fill_value=0)
# See https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.reindex.html

df_combined_daily_keypresses.index.name = 'Date'
df_combined_daily_keypresses.reset_index(inplace=True)
```

Now that I have a more complete daily keypress history, I can begin performing analyses on this data. First, I will use the rolling() function within Pandas to calculate 7-, 28-, and 365-day moving averages. Next, I will use the Series.rank() function twice to (1) calculate the percentile of each keypress and (2) determine each keypress's rank within the entire DataFrame.

```
[ ]: df_combined_daily_keypresses['7_day_ma'] = df_combined_daily_keypresses[
    'Keypresses'].rolling(7).mean()
df_combined_daily_keypresses['28_day_ma'] = df_combined_daily_keypresses[
    'Keypresses'].rolling(28).mean()
# I switched from a 30-day MA to a 28-day MA because my average keypresses vary
# significantly by weekday, meaning that a 30-day average would be skewed by
# the number of Saturdays and Sundays present in the data.
df_combined_daily_keypresses['365_day_ma'] = df_combined_daily_keypresses[
    'Keypresses'].rolling(365).mean()

df_combined_daily_keypresses['percentile'] = 100*df_combined_daily_keypresses[
    'Keypresses'].rank(pct=True)
# The pct=True argument generates percentile values for each keypress value.
df_combined_daily_keypresses['rank'] = df_combined_daily_keypresses[
    'Keypresses'].rank(ascending = False)
# Ascending = False instructs the function to assign the lowest number
# (e.g. 1) to the highest value.
# If two dates are tied, the rank may end in a 0.5. That's why these ranks
# appear as floats but not integers.
df_combined_daily_keypresses
```

```
[ ]:      Date  Keypresses    7_day_ma    28_day_ma    365_day_ma  \
0    2020-09-21      19081         NaN         NaN         NaN
1    2020-09-22      32771         NaN         NaN         NaN
2    2020-09-23      32065         NaN         NaN         NaN
3    2020-09-24      34698         NaN         NaN         NaN
4    2020-09-25      47038         NaN         NaN         NaN
...    ...    ...    ...    ...    ...
1268 2024-03-12      65458  32027.857143  19582.678571  25437.578082
1269 2024-03-13      57509  33075.571429  21636.392857  25499.835616
1270 2024-03-14      67124  42183.857143  23924.678571  25595.394521
1271 2024-03-15      36963  47218.142857  24798.500000  25619.561644
1272 2024-03-16      29865  45242.714286  24725.000000  25655.810959

      percentile  rank
0      40.298507  761.0
1      65.514533  440.0
2      64.179104  457.0
3      69.010212  395.5
4      87.274156  163.0
...    ...    ...
1268    97.800471   29.0
1269    94.658288   69.0
1270    98.350353   22.0
1271    73.448547  339.0
1272    59.858602  512.0

[1273 rows x 7 columns]
```

Next, I'll add in weekdays using the `Series.map()` function within Pandas.

```
[ ]: weekday_dict = {0: 'Monday', 1: 'Tuesday', 2: 'Wednesday', 3: 'Thursday', 4: 'Friday',
5: 'Saturday', 6: 'Sunday'}
# weekday numbers in Python begin with 0 for Monday and end with 6 for Sunday.
# See https://docs.python.org/3/library/datetime.html#datetime.date.weekday
weekday_dict
```

```
[ ]: {0: 'Monday',
1: 'Tuesday',
2: 'Wednesday',
3: 'Thursday',
4: 'Friday',
5: 'Saturday',
6: 'Sunday'}
```

The following cell adds a 'Weekday' column to the DataFrame by (1) calculating the numerical weekday values for each date in the 'Date' column, then (2) using `weekday_dict` to map these numerical values to the weekday names.

```
[ ]: df_combined_daily_keypresses['Weekday'] = df_combined_daily_keypresses[
    'Date'].dt.weekday.map(weekday_dict)
```

Here are my daily keypress statistics for the last 10 days. Note the presence of the moving average, percentile, rank, and weekday columns.

```
[ ]: df_combined_daily_keypresses.tail(10)
```

```
[ ]:
      Date  Keypresses    7_day_ma    28_day_ma    365_day_ma  \
1263 2024-03-07         3366  14614.142857  15419.678571  25332.695890
1264 2024-03-08         1723  14731.857143  15409.500000  25286.715068
1265 2024-03-09        43693  20970.142857  16807.500000  25317.641096
1266 2024-03-10          111  20977.142857  16149.178571  25273.569863
1267 2024-03-11        59669  28532.000000  18040.607143  25380.912329
1268 2024-03-12        65458  32027.857143  19582.678571  25437.578082
1269 2024-03-13        57509  33075.571429  21636.392857  25499.835616
1270 2024-03-14        67124  42183.857143  23924.678571  25595.394521
1271 2024-03-15        36963  47218.142857  24798.500000  25619.561644
1272 2024-03-16        29865  45242.714286  24725.000000  25655.810959
```

```

      percentile    rank  Weekday
1263   16.732129  1061.0  Thursday
1264   13.040063  1108.0    Friday
1265   83.582090   210.0  Saturday
1266    7.305577  1181.0    Sunday
1267   95.679497    56.0    Monday
1268   97.800471    29.0   Tuesday
1269   94.658288    69.0  Wednesday
1270   98.350353    22.0  Thursday
1271   73.448547   339.0    Friday
1272   59.858602   512.0  Saturday
```

2 Data Analysis

I'll start my data analysis by calculating some summary statistics. In the output below, 'count' shows the number of days since the earliest entry in my database, rather than the number of days for which I have keypress data (as the table also includes days without any keypresses).

```
[ ]: df_combined_daily_keypresses['Keypresses'].describe()
```

```
[ ]: count      1273.000000
      mean      25355.311862
      std       18766.514060
      min         0.000000
      25%       8712.000000
      50%      24829.000000
      75%      37989.000000
```

```
max      121833.000000
Name: Keypresses, dtype: float64
```

Next, I'll calculate my top 50 daily keypress totals. Note that the top-ranked date has a rank of 1 and a percentile of 100.

```
[ ]: df_max_keypresses = df_combined_daily_keypresses.sort_values(
      'Keypresses', ascending = False).reset_index(drop=True)
df_max_keypresses.insert(0, 'Rank', df_max_keypresses.index+1)
keypress_difference_list = [
    df_max_keypresses.iloc[i]['Keypresses'] - df_max_keypresses.iloc[i+1][
        'Keypresses'] for i in range(len(df_max_keypresses) -1 )]
# The above list comprehension calculates the difference
# between each row and the row below it. This isn't possible for the final row,
# so I'll instead append a np.NaN to it.)
keypress_difference_list.append(np.NaN)
df_max_keypresses['difference_from_lower_rank'] = keypress_difference_list
df_max_keypresses.head(50)
```

```
[ ]:   Rank      Date  Keypresses    7_day_ma    28_day_ma    365_day_ma  \
0      1  2024-01-08    121833  55373.857143  42174.785714  25960.498630
1      2  2023-11-18    104525  48377.000000  36835.035714  23188.216438
2      3  2024-01-03     94380  35844.142857  41919.392857  25558.646575
3      4  2022-12-05     90447  34728.714286  26518.821429  20476.994521
4      5  2023-12-07     81369  56303.857143  43351.821429  24036.515068
5      6  2023-12-26     79499  45022.571429  47711.428571  25165.372603
6      7  2024-01-02     78643  23509.571429  40957.500000  25406.241096
7      8  2021-02-22     77753  47983.000000  41390.857143           NaN
8      9  2020-11-23     76313  50377.571429  39924.107143           NaN
9     10  2020-11-19     75444  47285.571429  37945.892857           NaN
10     11  2021-11-09     73788  37843.285714  30462.071429  29584.802740
11     12  2023-11-29     73736  33759.571429  38933.714286  23717.972603
12     13  2022-05-03     72910  30511.428571  28489.285714  24401.624658
13     14  2023-12-19     71406  50155.142857  42303.535714  24531.665753
14     15  2021-02-08     70591  30177.428571  34437.357143           NaN
15     16  2024-01-24     68403  42259.428571  35836.928571  26463.991781
16     17  2021-04-06     68313  40455.857143  40719.178571           NaN
17     18  2021-02-15     67930  46174.571429  40268.535714           NaN
18     19  2020-11-30     67533  20747.428571  36843.714286           NaN
19     20  2021-02-17     67505  45946.571429  40801.607143           NaN
20     21  2023-12-06     67447  49626.857143  41540.107143  23935.002740
21     22  2024-03-14     67124  42183.857143  23924.678571  25595.394521
22     23  2021-10-12     66517  39971.285714  36729.892857  29886.202740
23     24  2021-10-14     65935  46708.142857  37694.500000  30001.038356
24     25  2021-09-23     65907  33891.142857  27215.107143  30006.980822
25     26  2023-11-28     65836  23391.000000  37445.035714  23572.430137
26     27  2023-12-02     65545  50889.857143  41582.607143  23857.506849
```


| | | | | | | |
|----|----|------------|-------|--------------|--------------|--------------|
| 27 | 28 | 2021-09-27 | 65464 | 35955.714286 | 29346.285714 | 29877.052055 |
| 28 | 29 | 2024-03-12 | 65458 | 32027.857143 | 19582.678571 | 25437.578082 |
| 29 | 30 | 2023-12-05 | 65076 | 50525.285714 | 41125.964286 | 23875.378082 |
| 30 | 31 | 2021-01-25 | 64903 | 43493.714286 | 27181.714286 | NaN |
| 31 | 32 | 2020-10-01 | 64873 | 41882.285714 | NaN | NaN |
| 32 | 33 | 2023-05-31 | 64683 | 26243.000000 | 21224.750000 | 20534.572603 |
| 33 | 34 | 2023-11-21 | 63528 | 58554.142857 | 40021.678571 | 23578.621918 |
| 34 | 35 | 2022-08-22 | 63149 | 23130.428571 | 21345.785714 | 22784.887671 |
| 35 | 36 | 2024-01-12 | 62973 | 45097.000000 | 40422.428571 | 26080.668493 |
| 36 | 37 | 2021-10-26 | 62922 | 30820.000000 | 37036.107143 | 29899.287671 |
| 37 | 38 | 2021-04-07 | 62668 | 41278.000000 | 40999.857143 | NaN |
| 38 | 39 | 2022-04-26 | 62596 | 31803.714286 | 29291.607143 | 23983.284932 |
| 39 | 40 | 2021-02-19 | 62473 | 47705.571429 | 41579.714286 | NaN |
| 40 | 41 | 2022-05-26 | 62030 | 21439.571429 | 20150.964286 | 23948.786301 |
| 41 | 42 | 2020-12-15 | 61969 | 44707.000000 | 35868.785714 | NaN |
| 42 | 43 | 2020-12-03 | 61437 | 36799.571429 | 36592.642857 | NaN |
| 43 | 44 | 2023-11-19 | 61303 | 51730.000000 | 39024.428571 | 23319.736986 |
| 44 | 45 | 2022-04-25 | 61200 | 28113.428571 | 28472.321429 | 23845.347945 |
| 45 | 46 | 2023-09-01 | 61187 | 35851.857143 | 28069.035714 | 21375.032877 |
| 46 | 47 | 2023-12-15 | 61156 | 44444.000000 | 46239.678571 | 24283.523288 |
| 47 | 48 | 2022-01-21 | 60953 | 38414.857143 | 17726.821429 | 27924.994521 |
| 48 | 49 | 2023-12-13 | 60893 | 44206.428571 | 45181.571429 | 24171.989041 |
| 49 | 50 | 2020-10-08 | 60890 | 42957.571429 | NaN | NaN |

| | percentile | rank | Weekday | difference_from_lower_rank |
|----|------------|------|-----------|----------------------------|
| 0 | 100.000000 | 1.0 | Monday | 17308.0 |
| 1 | 99.921445 | 2.0 | Saturday | 10145.0 |
| 2 | 99.842891 | 3.0 | Wednesday | 3933.0 |
| 3 | 99.764336 | 4.0 | Monday | 9078.0 |
| 4 | 99.685782 | 5.0 | Thursday | 1870.0 |
| 5 | 99.607227 | 6.0 | Tuesday | 856.0 |
| 6 | 99.528672 | 7.0 | Tuesday | 890.0 |
| 7 | 99.450118 | 8.0 | Monday | 1440.0 |
| 8 | 99.371563 | 9.0 | Monday | 869.0 |
| 9 | 99.293009 | 10.0 | Thursday | 1656.0 |
| 10 | 99.214454 | 11.0 | Tuesday | 52.0 |
| 11 | 99.135899 | 12.0 | Wednesday | 826.0 |
| 12 | 99.057345 | 13.0 | Tuesday | 1504.0 |
| 13 | 98.978790 | 14.0 | Tuesday | 815.0 |
| 14 | 98.900236 | 15.0 | Monday | 2188.0 |
| 15 | 98.821681 | 16.0 | Wednesday | 90.0 |
| 16 | 98.743126 | 17.0 | Tuesday | 383.0 |
| 17 | 98.664572 | 18.0 | Monday | 397.0 |
| 18 | 98.586017 | 19.0 | Monday | 28.0 |
| 19 | 98.507463 | 20.0 | Wednesday | 58.0 |
| 20 | 98.428908 | 21.0 | Wednesday | 323.0 |
| 21 | 98.350353 | 22.0 | Thursday | 607.0 |

| | | | | |
|----|-----------|------|-----------|--------|
| 22 | 98.271799 | 23.0 | Tuesday | 582.0 |
| 23 | 98.193244 | 24.0 | Thursday | 28.0 |
| 24 | 98.114690 | 25.0 | Thursday | 71.0 |
| 25 | 98.036135 | 26.0 | Tuesday | 291.0 |
| 26 | 97.957581 | 27.0 | Saturday | 81.0 |
| 27 | 97.879026 | 28.0 | Monday | 6.0 |
| 28 | 97.800471 | 29.0 | Tuesday | 382.0 |
| 29 | 97.721917 | 30.0 | Tuesday | 173.0 |
| 30 | 97.643362 | 31.0 | Monday | 30.0 |
| 31 | 97.564808 | 32.0 | Thursday | 190.0 |
| 32 | 97.486253 | 33.0 | Wednesday | 1155.0 |
| 33 | 97.407698 | 34.0 | Tuesday | 379.0 |
| 34 | 97.329144 | 35.0 | Monday | 176.0 |
| 35 | 97.250589 | 36.0 | Friday | 51.0 |
| 36 | 97.172035 | 37.0 | Tuesday | 254.0 |
| 37 | 97.093480 | 38.0 | Wednesday | 72.0 |
| 38 | 97.014925 | 39.0 | Tuesday | 123.0 |
| 39 | 96.936371 | 40.0 | Friday | 443.0 |
| 40 | 96.857816 | 41.0 | Thursday | 61.0 |
| 41 | 96.779262 | 42.0 | Tuesday | 532.0 |
| 42 | 96.700707 | 43.0 | Thursday | 134.0 |
| 43 | 96.622152 | 44.0 | Sunday | 103.0 |
| 44 | 96.543598 | 45.0 | Monday | 13.0 |
| 45 | 96.465043 | 46.0 | Friday | 31.0 |
| 46 | 96.386489 | 47.0 | Friday | 203.0 |
| 47 | 96.307934 | 48.0 | Friday | 60.0 |
| 48 | 96.229379 | 49.0 | Wednesday | 3.0 |
| 49 | 96.150825 | 50.0 | Thursday | 70.0 |

Next, I'll create a visualization of these 50 dates. I will use Plotly instead of Matplotlib so that I can create both interactive (.html) and static (.png) versions of each chart. The static versions are easier to share, but the interactive versions are easier to analyze, as you can hover over the chart to get more information and zoom in on areas of particular interest.

First, I'll create a function that will make it easier to save .html and .png versions of these charts:

```
[ ]: def save_chart(fig, file_name,
interactive_graphs_folder = interactive_graphs_folder,
static_graphs_folder = static_graphs_folder):
    '''Saves a Plotly figure as a .HTML and .PNG file.
    fig: The Plotly figure to save.
    file_name: The filename to use. Don't include the extension.'''
    fig.write_html(
        interactive_graphs_folder+'/'+file_name+'.html')
    # Saving the interactive chart to a .png file:
    fig.write_image(
        static_graphs_folder+'/'+file_name+'.png',
        width = default_image_width, height = default_image_height,
```

```
scale = default_image_scale)
```

```
[ ]: # Generating the interactive chart:
fig_max_keypresses = px.bar(df_max_keypresses.head(50),
x = 'Rank', y = 'Keypresses', title = 'Top 50 Daily Keypress Totals',
text = 'Keypresses')

save_chart(fig_max_keypresses, 'top_50_keypresses')
```

Here's a look at the chart:

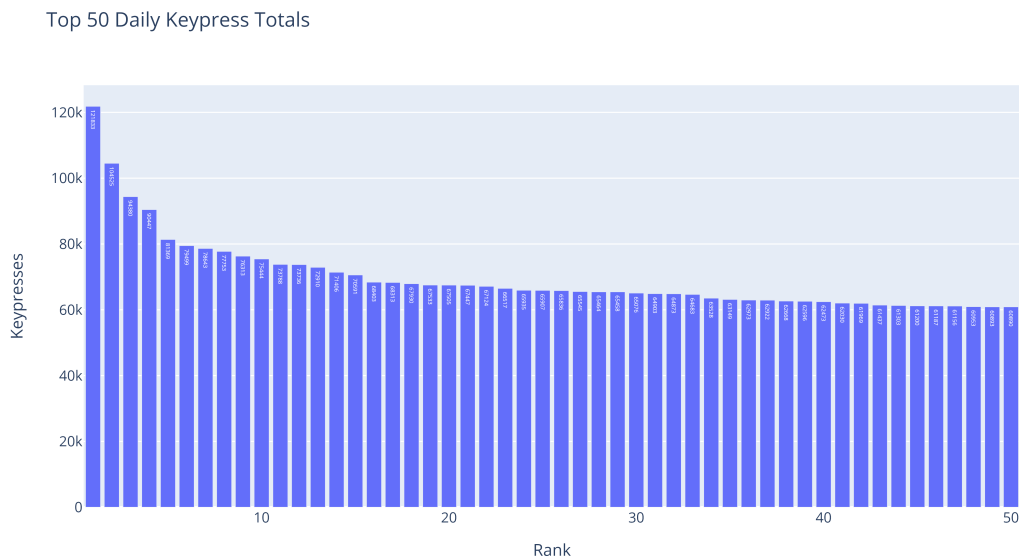
Note: this and other .html-based charts will probably not display for you within GitHub. In order to view them, you will need to download the files from GitHub (e.g. by cloning them) and view them on your computer.

```
[ ]: fig_max_keypresses
```

Here's a copy of the static chart in case the .html chart above did not load for you: (I'll also display static versions of later charts as well.)

```
[ ]: Image(static_graphs_folder+'top_50_keypresses.png')
```

```
[ ]:
```



2.1 Keypress percentile data:

First, I'll calculate the keypress totals equal to the 90th through 100th percentiles (in 1-percentile increments).

```
[ ]: df_combined_daily_keypresses['Keypresses'].describe(
    percentiles=np.linspace(0.9,1,11))[5:-1]
# The first four rows and final row provide additional descriptive statistics,
# so we can get rid of them
# by adding [4:-1] to the end of this line.
# There is probably a more elegant solution that retrieves only percentile
# data, but this option is fairly straightforward.
```

```
[ ]: 90%      50010.40
      91%      51265.52
      92%      53222.96
      93%      54594.20
      94%      56848.44
      95%      58115.00
      96%      60765.44
      97%      62576.32
      98%      65707.96
      99%      71827.12
     100%     121833.00
Name: Keypresses, dtype: float64
```

Scipy's `percentileofscore()` function can be used to calculate the percentile corresponding to a specific keypress total. For instance, here's the percentile for a day with only a single keypress: (The percentile may be higher than you'd expect due to the presence of many days with 0 keypresses)

```
[ ]: percentileofscore(df_combined_daily_keypresses['Keypresses'], 1)

[ ]: 5.813040062843676
```

Next, I'll create a DataFrame showing the keypresses corresponding to every 5th percentile.

```
[ ]: df_percentiles = df_combined_daily_keypresses['Keypresses'].describe(
    percentiles=np.linspace(0,1,21))[4:-1].reset_index().sort_values(
        'Keypresses', ascending = False).reset_index(drop=True).rename(
            columns={'index':'percentile'})
# Calling reset_index() twice appears inefficient, but it makes it easier
# to sort by a particular value (in this case, keypresses).
keypress_difference_list = [df_percentiles.iloc[
    i, 1] - df_percentiles.iloc[i+1, 1] for i in range(len(df_percentiles) -1 )]
# This list comprehension calculates the difference between each row and
# the row below it. This isn't possible for the final row,
# so we'll instead append a np.NaN to it.)
keypress_difference_list.append(np.NaN)
df_percentiles['difference_from_lower_percentile'] = keypress_difference_list
df_percentiles
```

```
[ ]: percentile  Keypresses  difference_from_lower_percentile
0          100%    121833.0                63718.0
1           95%     58115.0                 8104.6
2           90%     50010.4                 5066.0
3           85%     44944.4                 3762.6
4           80%     41181.8                 3192.8
5           75%     37989.0                 2817.4
6           70%     35171.6                 2901.2
7           65%     32270.4                 2344.0
8           60%     29926.4                 2958.2
9           55%     26968.2                 2139.2
10          50%     24829.0                 2930.2
11          45%     21898.8                 2929.0
12          40%     18969.8                 2792.8
13          35%     16177.0                 3502.2
14          30%     12674.8                 3962.8
15          25%      8712.0                 2770.2
16          20%      5941.8                 3448.8
17          15%      2493.0                 1716.6
18          10%       776.4                  776.4
19           5%         0.0                   0.0
20           0%         0.0                  NaN
```

The following chart shows the difference in keypresses between the different percentiles on this list.

```
[ ]: px.bar(df_percentiles, x = 'percentile',
           y = 'difference_from_lower_percentile')
```

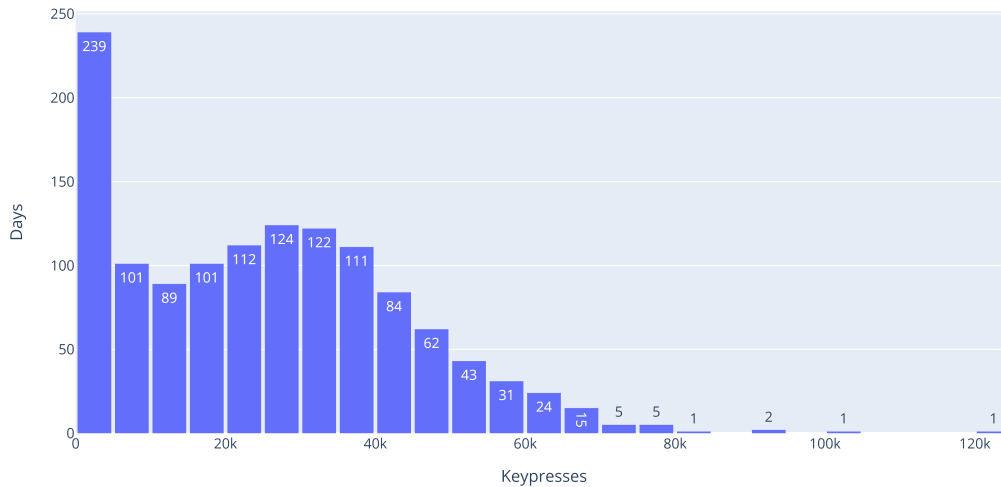
And here's a histogram that shows the frequency of different keypress ranges:

```
[ ]: fig_keypress_hist = px.histogram(df_combined_daily_keypresses,
x = 'Keypresses', nbins = 40, text_auto = True,
title = 'Histogram of Daily Keypress Totals')
# See https://plotly.com/python/histograms/
fig_keypress_hist.update_layout(bargap=0.1, yaxis_title = 'Days')
# https://plotly.com/python/histograms/
save_chart(fig_keypress_hist, 'daily_keypress_histogram')
fig_keypress_hist
```

```
[ ]: Image(static_graphs_folder+'daily_keypress_histogram.png')
```

```
[ ]:
```

Histogram of Daily Keypress Totals



I'll now create a pivot table that shows my average keypresses per weekday:

```
[ ]: df_weekday_pivot = df_combined_daily_keypresses.pivot_table(
    index = 'Weekday', values = 'Keypresses', aggfunc = 'mean').sort_values(
        'Weekday', ascending = False).reset_index()
# Adding in a Weekday_Order column so that
# weekdays can be sorted chronologically:
df_weekday_pivot['Weekday_Order'] = df_weekday_pivot['Weekday'].map(
    {'Sunday':0, 'Monday':1, 'Tuesday':2, 'Wednesday':3,
    'Thursday':4, 'Friday':5, 'Saturday':6})
df_weekday_pivot
```

```
[ ]:      Weekday  Keypresses  Weekday_Order
0  Wednesday    31415.642857             3
1   Tuesday     32974.538462             2
2  Thursday     30795.065934             4
3    Sunday     11635.508287             0
4  Saturday     14914.456044             6
5   Monday      31004.730769             1
6   Friday      24671.857143             5
```

```
[ ]: df_weekday_pivot.sort_values('Weekday_Order', inplace = True)
df_weekday_pivot
```

```
[ ]:      Weekday  Keypresses  Weekday_Order
3    Sunday     11635.508287             0
5   Monday      31004.730769             1
```

| | | | |
|---|-----------|--------------|---|
| 1 | Tuesday | 32974.538462 | 2 |
| 0 | Wednesday | 31415.642857 | 3 |
| 2 | Thursday | 30795.065934 | 4 |
| 6 | Friday | 24671.857143 | 5 |
| 4 | Saturday | 14914.456044 | 6 |

Graphing my average keypresses per weekday:

```
[ ]: fig_keypresses_by_weekday = px.bar(df_weekday_pivot, x = 'Weekday',
y = 'Keypresses', text_auto = '.0f', color = 'Keypresses',
color_continuous_scale = 'PrGn',
title = 'Average Keypresses by Weekday')
# For text_auto, see:
# https://plotly.com/python-api-reference/generated/plotly.express.bar


# For color scale options and usage,
# visit https://plotly.com/python/builtin-colorscales/
fig_keypresses_by_weekday.update_traces(marker_line_color='rgb(0,0,0)',
marker_line_width=0.5, opacity=1)

fig_keypresses_by_weekday

save_chart(fig_keypresses_by_weekday, file_name = 'keypresses_by_weekday')
```

```
[ ]: Image(static_graphs_folder+'keypresses_by_weekday.png')
```

```
[ ]:
```



2.1.1 Total keypresses since first date in DataFrame:

```
[ ]: print("Total keypresses since", str(
df_combined_daily_keypresses.iloc[0]['Date'])+":", '{:,.}'.format(
sum(df_combined_daily_keypresses['Keypresses'])))
```

Total keypresses since 2020-09-21 00:00:00: 32,277,312

2.1.2 Keypresses over the past 50 days:

```
[ ]: df_combined_daily_keypresses.tail(50) # Last 50 days
```

```
[ ]:
      Date  Keypresses    7_day_ma    28_day_ma    365_day_ma  \
1223 2024-01-27         0  22625.142857  34161.964286  26284.336986
1224 2024-01-28       5283  22415.714286  33982.107143  26279.671233
1225 2024-01-29      8541  18517.285714  33549.000000  26283.747945
1226 2024-01-30      3336  12223.285714  30859.464286  26184.572603
1227 2024-01-31      2201   2765.857143  27567.357143  26117.723288
1228 2024-02-01         0   2765.857143  27149.500000  26009.164384
1229 2024-02-02     13198   4651.285714  25718.678571  26029.260274
1230 2024-02-03     11912   6353.000000  25658.892857  26028.054795
1231 2024-02-04     23617   8972.142857  25994.714286  26072.178082
1232 2024-02-05       1250   7930.571429  21688.178571  26073.361644
1233 2024-02-06        769   7563.857143  20954.178571  26034.246575
1234 2024-02-07       3313   7722.714286  19823.321429  25960.608219
1235 2024-02-08       8530   8941.285714  18457.428571  25908.405479
1236 2024-02-09       2008   7342.714286  16280.107143  25819.479452
1237 2024-02-10      4549   6290.857143  16216.678571  25750.994521
1238 2024-02-11     18544   5566.142857  16719.428571  25799.213699
1239 2024-02-12       6709   6346.000000  16628.285714  25808.572603
1240 2024-02-13     22280   9419.000000  16351.607143  25819.561644
1241 2024-02-14         5   8946.428571  15423.607143  25781.564384
1242 2024-02-15      3052   8163.857143  13496.000000  25746.882192
1243 2024-02-16     12496   9662.142857  12791.071429  25670.002740
1244 2024-02-17     31923  13572.714286  12210.428571  25624.249315
1245 2024-02-18       2512  11282.428571  12059.107143  25579.504110
1246 2024-02-19     13534  12257.428571  11262.821429  25611.495890
1247 2024-02-20     13287  10972.714286  10044.714286  25642.627397
1248 2024-02-21     43274  17154.000000   9147.250000  25662.504110
1249 2024-02-22     16772  19114.000000   9746.250000  25604.808219
1250 2024-02-23     31806  21872.571429  10882.178571  25557.430137
1251 2024-02-24     19710  20127.857143  11586.107143  25519.690411
1252 2024-02-25         0  19769.000000  11397.428571  25491.156164
1253 2024-02-26      2595  18206.285714  11185.071429  25498.265753
1254 2024-02-27      1567  16532.000000  11121.892857  25431.328767
1255 2024-02-28     57119  18509.857143  13083.250000  25531.758904
1256 2024-02-29     25710  19786.714286  14001.464286  25505.895890
1257 2024-03-01       899  15371.428571  13562.214286  25386.567123
```


| | | | | | |
|------|------------|-------|--------------|--------------|--------------|
| 1258 | 2024-03-02 | 25 | 12559.285714 | 13137.678571 | 25382.693151 |
| 1259 | 2024-03-03 | 62 | 12568.142857 | 12296.428571 | 25382.863014 |
| 1260 | 2024-03-04 | 6785 | 13166.714286 | 12494.107143 | 25355.764384 |
| 1261 | 2024-03-05 | 40987 | 18798.142857 | 13930.464286 | 25359.684932 |
| 1262 | 2024-03-06 | 50175 | 17806.142857 | 15604.107143 | 25417.536986 |
| 1263 | 2024-03-07 | 3366 | 14614.142857 | 15419.678571 | 25332.695890 |
| 1264 | 2024-03-08 | 1723 | 14731.857143 | 15409.500000 | 25286.715068 |
| 1265 | 2024-03-09 | 43693 | 20970.142857 | 16807.500000 | 25317.641096 |
| 1266 | 2024-03-10 | 111 | 20977.142857 | 16149.178571 | 25273.569863 |
| 1267 | 2024-03-11 | 59669 | 28532.000000 | 18040.607143 | 25380.912329 |
| 1268 | 2024-03-12 | 65458 | 32027.857143 | 19582.678571 | 25437.578082 |
| 1269 | 2024-03-13 | 57509 | 33075.571429 | 21636.392857 | 25499.835616 |
| 1270 | 2024-03-14 | 67124 | 42183.857143 | 23924.678571 | 25595.394521 |
| 1271 | 2024-03-15 | 36963 | 47218.142857 | 24798.500000 | 25619.561644 |
| 1272 | 2024-03-16 | 29865 | 45242.714286 | 24725.000000 | 25655.810959 |

| | percentile | rank | Weekday |
|------|------------|--------|-----------|
| 1223 | 2.906520 | 1237.0 | Saturday |
| 1224 | 19.167321 | 1030.0 | Sunday |
| 1225 | 24.587588 | 961.0 | Monday |
| 1226 | 16.653574 | 1062.0 | Tuesday |
| 1227 | 14.296936 | 1092.0 | Wednesday |
| 1228 | 2.906520 | 1237.0 | Thursday |
| 1229 | 31.029065 | 879.0 | Friday |
| 1230 | 28.986646 | 905.0 | Saturday |
| 1231 | 47.918303 | 664.0 | Sunday |
| 1232 | 11.861744 | 1123.0 | Monday |
| 1233 | 10.054988 | 1146.0 | Tuesday |
| 1234 | 16.575020 | 1063.0 | Wednesday |
| 1235 | 24.509034 | 962.0 | Thursday |
| 1236 | 13.904163 | 1097.0 | Friday |
| 1237 | 18.224666 | 1042.0 | Saturday |
| 1238 | 39.120189 | 776.0 | Sunday |
| 1239 | 21.681068 | 998.0 | Monday |
| 1240 | 45.561665 | 694.0 | Tuesday |
| 1241 | 5.970149 | 1198.0 | Wednesday |
| 1242 | 16.025137 | 1070.0 | Thursday |
| 1243 | 29.693637 | 896.0 | Friday |
| 1244 | 63.668500 | 463.5 | Saturday |
| 1245 | 15.082482 | 1082.0 | Sunday |
| 1246 | 31.736057 | 870.0 | Monday |
| 1247 | 31.186174 | 877.0 | Tuesday |
| 1248 | 83.032207 | 217.0 | Wednesday |
| 1249 | 36.135114 | 814.0 | Thursday |
| 1250 | 63.393559 | 467.0 | Friday |
| 1251 | 41.005499 | 752.0 | Saturday |
| 1252 | 2.906520 | 1237.0 | Sunday |

| | | | |
|------|-----------|--------|-----------|
| 1253 | 15.475255 | 1077.0 | Monday |
| 1254 | 12.725844 | 1112.0 | Tuesday |
| 1255 | 94.344069 | 73.0 | Wednesday |
| 1256 | 51.610369 | 617.0 | Thursday |
| 1257 | 10.526316 | 1140.0 | Friday |
| 1258 | 6.205813 | 1195.0 | Saturday |
| 1259 | 6.755695 | 1188.0 | Sunday |
| 1260 | 21.838178 | 996.0 | Monday |
| 1261 | 79.497251 | 262.0 | Tuesday |
| 1262 | 90.259230 | 125.0 | Wednesday |
| 1263 | 16.732129 | 1061.0 | Thursday |
| 1264 | 13.040063 | 1108.0 | Friday |
| 1265 | 83.582090 | 210.0 | Saturday |
| 1266 | 7.305577 | 1181.0 | Sunday |
| 1267 | 95.679497 | 56.0 | Monday |
| 1268 | 97.800471 | 29.0 | Tuesday |
| 1269 | 94.658288 | 69.0 | Wednesday |
| 1270 | 98.350353 | 22.0 | Thursday |
| 1271 | 73.448547 | 339.0 | Friday |
| 1272 | 59.858602 | 512.0 | Saturday |

The following cell outputs various keypress statistics. When this script is run hourly, these statistics (along with other ones) will then appear in the console terminal.

```
[ ]: days_with_data = len(df_combined_daily_keypresses)
# The following column cell shows the ranks immediately above the ranks for
↳ the most recent day.
keypresses_today = df_combined_daily_keypresses.iloc[-1]['Keypresses']
percentile_today = df_combined_daily_keypresses.iloc[-1]['percentile']
rank_today = df_combined_daily_keypresses.iloc[-1]['rank']
print("Ranks are out of", days_with_data, "days.")
print(f"Today's keypresses: {keypresses_today}")
print(f"Your keypress totals yesterday and 7, 28, and 365 days ago were \
{df_combined_daily_keypresses.iloc[-2]['Keypresses']}, \
{df_combined_daily_keypresses.iloc[-8]['Keypresses']}, \
{df_combined_daily_keypresses.iloc[-29]['Keypresses']}, \
and {df_combined_daily_keypresses.iloc[-366]['Keypresses']}, respectively.")
# If your keypresses today are higher than these values, the moving averages
# associated with those values will increase.
print(f"Today's percentile: {round(percentile_today, 3)}")
print(f"Today's rank: {rank_today} \
(in front of {days_with_data - rank_today} days)")
```

Ranks are out of 1273 days.

Today's keypresses: 29865

Your keypress totals yesterday and 7, 28, and 365 days ago were 36963, 43693, 31923, and 16634, respectively.

Today's percentile: 59.859

Today's rank: 512.0 (in front of 761.0 days)

Days ranked just ahead of today (along with today's rank):

```
[ ]: df_days_with_higher_keypresses = df_combined_daily_keypresses.sort_values(
    'rank').query("rank <= @rank_today").tail(11)
keypress_difference_list = [df_days_with_higher_keypresses.iloc[i][
    'Keypresses'] - df_days_with_higher_keypresses.iloc[i+1][
    'Keypresses'] for i in range(len(df_days_with_higher_keypresses) -1 )]
keypress_difference_list.append(np.NaN)
df_days_with_higher_keypresses[
    'diff_from_following_day'] = keypress_difference_list
df_days_with_higher_keypresses[
    'diff_from_current_day'] = df_days_with_higher_keypresses[
    'Keypresses'] - df_days_with_higher_keypresses.iloc[-1]['Keypresses']
df_days_with_higher_keypresses
```

```
[ ]:      Date  Keypresses    7_day_ma    28_day_ma    365_day_ma  \
1026 2023-07-14      30171  18003.857143  16546.821429  20932.805479
547  2022-03-22      30165  10191.857143  11713.392857  24423.895890
578  2022-04-22      30101  24341.857143  29521.500000  23799.232877
1173 2023-12-08      30078  52329.428571  44081.750000  23996.997260
1212 2024-01-16      30027  27829.285714  35886.607143  26081.838356
343  2021-08-30      30002  13560.857143  22669.000000          NaN
436  2021-12-01      29954  21674.571429  27999.821429  28964.024658
1148 2023-11-13      29944  31486.285714  27291.928571  22814.586301
120  2021-01-19      29922  22547.285714  18311.178571          NaN
737  2022-09-28      29908  18806.000000  11800.107143  21618.230137
1272 2024-03-16      29865  45242.714286  24725.000000  25655.810959
```

```
      percentile  rank  Weekday  diff_from_following_day  \
1026   60.644148  502.0   Friday              6.0
547    60.565593  503.0  Tuesday             64.0
578    60.487038  504.0   Friday             23.0
1173    60.408484  505.0   Friday             51.0
1212    60.329929  506.0  Tuesday             25.0
343     60.251375  507.0   Monday             48.0
436     60.172820  508.0 Wednesday             10.0
1148    60.094266  509.0   Monday             22.0
120     60.015711  510.0  Tuesday             14.0
737     59.937156  511.0 Wednesday             43.0
1272    59.858602  512.0 Saturday             NaN
```

```
      diff_from_current_day
1026                    306
547                     300
578                     236
1173                    213
```

| | |
|------|-----|
| 1212 | 162 |
| 343 | 137 |
| 436 | 89 |
| 1148 | 79 |
| 120 | 57 |
| 737 | 43 |
| 1272 | 0 |

Looking for days with identical non-zero keypress totals:

```
[ ]: duplicated_keypress_dates = df_combined_daily_keypresses[
df_combined_daily_keypresses.duplicated(
subset = 'Keypresses', keep = False)].query('Keypresses > 0').sort_values(
'Keypresses', ascending = False)
len(duplicated_keypress_dates)
```

```
[ ]: 24
```

```
[ ]: duplicated_keypress_dates
```

```
[ ]:
      Date  Keypresses    7_day_ma    28_day_ma    365_day_ma  \
32  2020-10-23      48908  41125.571429  37996.178571      NaN
183 2021-03-23      48908  36708.000000  32516.464286      NaN
315 2021-08-02      42704  30582.285714  29370.714286      NaN
1071 2023-08-28      42704  28179.142857  24481.535714  21111.791781
1218 2024-01-22      35830  33719.285714  34827.642857  26308.213699
746 2022-10-07      35830  22178.714286  14628.250000  21042.265753
3   2020-09-24      34698           NaN           NaN           NaN
124 2021-01-23      34698  34726.428571  24214.178571      NaN
318 2021-08-05      32184  29066.857143  29787.500000      NaN
676 2022-07-29      32184  26384.571429  26299.000000  23180.561644
1244 2024-02-17      31923  13572.714286  12210.428571  25624.249315
525 2022-02-28      31923  18683.142857  17425.357143  25764.104110
687 2022-08-09      26905  21784.571429  25372.785714  22972.580822
765 2022-10-26      26905  25745.714286  20627.357143  20359.476712
645 2022-06-28      26550  12407.857143  12670.821429  23373.895890
55  2020-11-15      26550  36325.285714  35269.428571      NaN
662 2022-07-15      25084  28583.714286  18304.428571  23319.049315
841 2023-01-10      25084  16131.571429  16824.607143  21120.632877
620 2022-06-03      23824  18572.857143  19212.571429  23844.123288
1050 2023-08-07      23824  17324.000000  23705.250000  20797.142466
321 2021-08-08       6318  30379.714286  30359.035714      NaN
632 2022-06-15       6318   6218.571429  15957.071429  23702.145205
538 2022-03-13        149   8908.285714  15542.392857  25160.463014
536 2022-03-11        149  10365.285714  15863.214286  25284.586301
```

| percentile | rank | Weekday |
|------------|------|---------|
|------------|------|---------|

| | | | |
|------|-----------|--------|-----------|
| 32 | 89.198743 | 138.5 | Friday |
| 183 | 89.198743 | 138.5 | Tuesday |
| 315 | 82.364493 | 225.5 | Monday |
| 1071 | 82.364493 | 225.5 | Monday |
| 1218 | 71.052632 | 369.5 | Monday |
| 746 | 71.052632 | 369.5 | Friday |
| 3 | 69.010212 | 395.5 | Thursday |
| 124 | 69.010212 | 395.5 | Saturday |
| 318 | 64.532600 | 452.5 | Thursday |
| 676 | 64.532600 | 452.5 | Friday |
| 1244 | 63.668500 | 463.5 | Saturday |
| 525 | 63.668500 | 463.5 | Monday |
| 687 | 54.791830 | 576.5 | Tuesday |
| 765 | 54.791830 | 576.5 | Wednesday |
| 645 | 53.927730 | 587.5 | Tuesday |
| 55 | 53.927730 | 587.5 | Sunday |
| 662 | 50.628437 | 629.5 | Friday |
| 841 | 50.628437 | 629.5 | Tuesday |
| 620 | 48.428908 | 657.5 | Friday |
| 1050 | 48.428908 | 657.5 | Monday |
| 321 | 21.013354 | 1006.5 | Sunday |
| 632 | 21.013354 | 1006.5 | Wednesday |
| 538 | 7.580518 | 1177.5 | Sunday |
| 536 | 7.580518 | 1177.5 | Friday |

```
[ ]: df_combined_daily_keypresses
```

```
[ ]:
      Date  Keypresses    7_day_ma    28_day_ma    365_day_ma  \
0  2020-09-21      19081         NaN         NaN         NaN
1  2020-09-22      32771         NaN         NaN         NaN
2  2020-09-23      32065         NaN         NaN         NaN
3  2020-09-24      34698         NaN         NaN         NaN
4  2020-09-25      47038         NaN         NaN         NaN
...
1268 2024-03-12      65458  32027.857143  19582.678571  25437.578082
1269 2024-03-13      57509  33075.571429  21636.392857  25499.835616
1270 2024-03-14      67124  42183.857143  23924.678571  25595.394521
1271 2024-03-15      36963  47218.142857  24798.500000  25619.561644
1272 2024-03-16      29865  45242.714286  24725.000000  25655.810959

      percentile  rank  Weekday
0      40.298507  761.0  Monday
1      65.514533  440.0  Tuesday
2      64.179104  457.0  Wednesday
3      69.010212  395.5  Thursday
4      87.274156  163.0  Friday
...           ...  ...      ...
```

| | | | |
|------|-----------|-------|-----------|
| 1268 | 97.800471 | 29.0 | Tuesday |
| 1269 | 94.658288 | 69.0 | Wednesday |
| 1270 | 98.350353 | 22.0 | Thursday |
| 1271 | 73.448547 | 339.0 | Friday |
| 1272 | 59.858602 | 512.0 | Saturday |

[1273 rows x 8 columns]

```
[ ]: df_combined_daily_keypresses['Date']
```

```
[ ]: 0      2020-09-21
      1      2020-09-22
      2      2020-09-23
      3      2020-09-24
      4      2020-09-25
      ...
     1268    2024-03-12
     1269    2024-03-13
     1270    2024-03-14
     1271    2024-03-15
     1272    2024-03-16
     Name: Date, Length: 1273, dtype: datetime64[ns]
```

2.2 Plotting Chronological Keypress Data

Now that we have a DataFrame showing daily keypresses and multiple moving averages, it's time to visualize it! The advantage of creating this chart within Plotly is that, being HTML-based, it is interactive in nature. Thus, you can hover over the lines to view the values corresponding to those lines and zoom in to get a closer look at a particular section of the graph. As before, though, this graph can also be saved as a static image.

```
[ ]: fig_keypresses_line_chart = px.line(df_combined_daily_keypresses, x = 'Date',
      y = ['Keypresses', '7_day_ma', '28_day_ma', '365_day_ma'],
      labels = {'variable': 'Metric', 'value': 'Keypresses'}, title =
      'Daily Keypresses and 7/28/365-day Moving Averages')
      # Note that multiple y values can be passed to the line chart.

      save_chart(fig_keypresses_line_chart, 'px_daily_keypresses_and_mas')


      fig_keypresses_line_chart
```

```
c:\Users\kburc\miniforge3\envs\ga15pyd\lib\site-
packages\_plotly_utils\basevalidators.py:105: FutureWarning:
```

The behavior of `DatetimeProperties.to_pydatetime` is deprecated, in a future version this will return a Series containing python datetime objects instead of an ndarray. To retain the old behavior, call ``np.array`` on the result

```
[ ]: Image(static_graphs_folder+'px_daily_keypresses_and_mas.png')
```

```
[ ]:
```



2.3 Monthly keypress totals:

```
[ ]: df_monthly_keypresses = df_combined_daily_keypresses.copy().set_index(
    'Date').resample('M').sum()['Keypresses'].reset_index()
df_monthly_keypresses['Month'] = df_monthly_keypresses['Date'].dt.to_period('M')
df_monthly_keypresses['Year'] = df_monthly_keypresses['Date'].dt.to_period('Y')
df_monthly_keypresses.drop('Date', axis = 1, inplace = True)
df_monthly_keypresses['Keypresses'] = df_monthly_keypresses.pop('Keypresses')
df_monthly_keypresses
```

C:\Users\kburc\AppData\Local\Temp\ipykernel_31088\3726139356.py:2:

FutureWarning:

'M' is deprecated and will be removed in a future version, please use 'ME' instead.

```
[ ]:
      Month  Year  Keypresses
0  2020-09  2020    346918
1  2020-10  2020   1140421
2  2020-11  2020   1061614
3  2020-12  2020    823763
4  2021-01  2021    980116
5  2021-02  2021   1088404
```

| | | | |
|----|---------|------|---------|
| 6 | 2021-03 | 2021 | 1142850 |
| 7 | 2021-04 | 2021 | 833464 |
| 8 | 2021-05 | 2021 | 700155 |
| 9 | 2021-06 | 2021 | 621894 |
| 10 | 2021-07 | 2021 | 881137 |
| 11 | 2021-08 | 2021 | 722970 |
| 12 | 2021-09 | 2021 | 919571 |
| 13 | 2021-10 | 2021 | 1051663 |
| 14 | 2021-11 | 2021 | 797115 |
| 15 | 2021-12 | 2021 | 515936 |
| 16 | 2022-01 | 2022 | 729233 |
| 17 | 2022-02 | 2022 | 487910 |
| 18 | 2022-03 | 2022 | 552559 |
| 19 | 2022-04 | 2022 | 834250 |
| 20 | 2022-05 | 2022 | 627263 |
| 21 | 2022-06 | 2022 | 412725 |
| 22 | 2022-07 | 2022 | 758759 |
| 23 | 2022-08 | 2022 | 691636 |
| 24 | 2022-09 | 2022 | 330566 |
| 25 | 2022-10 | 2022 | 706081 |
| 26 | 2022-11 | 2022 | 756509 |
| 27 | 2022-12 | 2022 | 720251 |
| 28 | 2023-01 | 2023 | 652155 |
| 29 | 2023-02 | 2023 | 580208 |
| 30 | 2023-03 | 2023 | 892786 |
| 31 | 2023-04 | 2023 | 327674 |
| 32 | 2023-05 | 2023 | 665769 |
| 33 | 2023-06 | 2023 | 628301 |
| 34 | 2023-07 | 2023 | 694009 |
| 35 | 2023-08 | 2023 | 809682 |
| 36 | 2023-09 | 2023 | 803639 |
| 37 | 2023-10 | 2023 | 715687 |
| 38 | 2023-11 | 2023 | 1156827 |
| 39 | 2023-12 | 2023 | 1292810 |
| 40 | 2024-01 | 2024 | 965577 |
| 41 | 2024-02 | 2024 | 392041 |
| 42 | 2024-03 | 2024 | 464414 |

3 Saving the updated version of this DataFrame to a .csv file:

```
[ ]: df_combined_daily_keypresses.to_csv(
      'data/df_combined_daily_keypresses_updated.csv')
```


3.1 Hourly keypress stats:

In order to calculate hourly keypress statistics, we'll need to create a new DataFrame that aggregates keypresses by hour instead of by day.

```
[ ]: hourly_keypress_db_list = []

for db_path in database_paths_list:
    hourly_keypress_db_list.append(generate_keypress_totals(
        database_path = db_path, level = 'hourly'))

df_hourly_keypresses = pd.concat([
    df for df in hourly_keypress_db_list]).reset_index(drop=True)

# As with my daily keypresses DataFrame, I'll use pivot_table() to group
# multiple rows for the same day and hour into a single row. (These multiple
# rows are the result of my using multiple computers during the same hour.)
df_hourly_keypresses = df_hourly_keypresses.pivot_table(index = [
    'Day', 'Hour'], values = 'Keypresses', aggfunc = 'sum').reset_index().
    ↪sort_values(['Day', 'Hour'])

df_hourly_keypresses.to_csv(
    'data/df_combined_hourly_keypresses.csv', index = False)
df_hourly_keypresses
```

```
[ ]:
      Day  Hour  Keypresses
0    2020-09-21    15         278
1    2020-09-21    16         567
2    2020-09-21    17        1218
3    2020-09-21    18        3300
4    2020-09-21    19        3888
...      ...  ...
14177 2024-03-16    13         862
14178 2024-03-16    14        6321
14179 2024-03-16    15        4146
14180 2024-03-16    16         494
14181 2024-03-16    17        1984
```

[14182 rows x 3 columns]

Recreating the DataFrame from a .csv file so that the following cells can be run by users who don't yet have their own WhatPulse database:

```
[ ]: df_hourly_keypresses = pd.read_csv('data/df_combined_hourly_keypresses.csv')
```

```
[ ]: df_hourly_keypresses['Day'] = pd.to_datetime(df_hourly_keypresses['Day'])
# Creating a combined day/hour column:
df_hourly_keypresses['Day_and_Hour'] = df_hourly_keypresses[
```

```

    'Day'] + pd.to_timedelta(df_hourly_keypresses['Hour'], unit = 'H')
df_hourly_keypresses.set_index('Day_and_Hour', inplace = True)
df_hourly_keypresses

```

C:\Users\kburc\AppData\Local\Temp\ipykernel_31088\137757146.py:4: FutureWarning:

'H' is deprecated and will be removed in a future version. Please use 'h' instead of 'H'.

```

[ ]:

```

| | Day | Hour | Keypresses |
|---------------------|------------|------|------------|
| Day_and_Hour | | | |
| 2020-09-21 15:00:00 | 2020-09-21 | 15 | 278 |
| 2020-09-21 16:00:00 | 2020-09-21 | 16 | 567 |
| 2020-09-21 17:00:00 | 2020-09-21 | 17 | 1218 |
| 2020-09-21 18:00:00 | 2020-09-21 | 18 | 3300 |
| 2020-09-21 19:00:00 | 2020-09-21 | 19 | 3888 |
| ... | ... | ... | ... |
| 2024-03-16 13:00:00 | 2024-03-16 | 13 | 862 |
| 2024-03-16 14:00:00 | 2024-03-16 | 14 | 6321 |
| 2024-03-16 15:00:00 | 2024-03-16 | 15 | 4146 |
| 2024-03-16 16:00:00 | 2024-03-16 | 16 | 494 |
| 2024-03-16 17:00:00 | 2024-03-16 | 17 | 1984 |

[14182 rows x 3 columns]

The following cells add in hours with 0 keypresses (of which there are many!). In order to include current/previous hours for the current day in my results, I'll add in keypresses up to the start of the next day (i.e. midnight), then limit the results so that they don't extend beyond the current hour.

Calculating tomorrow's date:

```

[ ]: last_date_for_hourly_keypress_log = last_date + datetime.timedelta(days = 1)
    last_date_for_hourly_keypress_log

```

```

[ ]: Timestamp('2024-03-17 00:00:00')

```

```

[ ]: pd.Timestamp.now()

```

```

[ ]: Timestamp('2024-03-16 17:44:35.253021')

```

Adding hours without keypresses to the DataFrame:

```

[ ]: full_hourly_date_range = pd.date_range(start = first_date,
    end = last_date_for_hourly_keypress_log, freq = 'h')
    df_hourly_keypresses = df_hourly_keypresses.reindex(full_hourly_date_range).
    ↪copy()

```

```
df_hourly_keypresses['Keypresses'] = df_hourly_keypresses['Keypresses'].
    ↪fillna(0)
df_hourly_keypresses['Keypresses'] = df_hourly_keypresses[
    'Keypresses'].astype('int')
# Retrieving date and hour values from the index:
df_hourly_keypresses['Day'] = df_hourly_keypresses.index.date
df_hourly_keypresses['Hour'] = df_hourly_keypresses.index.hour
df_hourly_keypresses
```

```
[ ]:
      Day  Hour  Keypresses
2020-09-21 00:00:00 2020-09-21    0      0
2020-09-21 01:00:00 2020-09-21    1      0
2020-09-21 02:00:00 2020-09-21    2      0
2020-09-21 03:00:00 2020-09-21    3      0
2020-09-21 04:00:00 2020-09-21    4      0
...
2024-03-16 20:00:00 2024-03-16   20      0
2024-03-16 21:00:00 2024-03-16   21      0
2024-03-16 22:00:00 2024-03-16   22      0
2024-03-16 23:00:00 2024-03-16   23      0
2024-03-17 00:00:00 2024-03-17    0      0
```

[30553 rows x 3 columns]

Limiting the results to the period leading up to the current hour:

```
[ ]: df_hourly_keypresses = df_hourly_keypresses[
    df_hourly_keypresses.index < pd.Timestamp.now()].copy()
df_hourly_keypresses.reset_index(drop=True,inplace=True)
```

Calculating rolling 24-hour keypress totals:

```
[ ]: df_hourly_keypresses['keypresses_over_last_24_hours'] = df_hourly_keypresses[
    'Keypresses'].rolling(24).sum()
df_hourly_keypresses
```

```
[ ]:
      Day  Hour  Keypresses  keypresses_over_last_24_hours
0    2020-09-21    0      0      NaN
1    2020-09-21    1      0      NaN
2    2020-09-21    2      0      NaN
3    2020-09-21    3      0      NaN
4    2020-09-21    4      0      NaN
...
30541 2024-03-16   13     862    30418.0
30542 2024-03-16   14    6321    34399.0
30543 2024-03-16   15    4146    34674.0
30544 2024-03-16   16     494    31006.0
30545 2024-03-16   17    1984    29926.0
```

[30546 rows x 4 columns]

Printing out recent hours with keypresses: (This data will also appear on the terminal window when the program is run automatically, which allows you to track your recent productivity.)

```
[ ]: print("Keypresses over the last 25 hours (excluding hours \
with 0 keypresses):\n",df_hourly_keypresses.iloc[-25:].query("Keypresses > 0"))
# Hours with 0 keypresses are removed in order to give the console output more
# space to fit on a single line.
```

Keypresses over the last 25 hours (excluding hours with 0 keypresses):

| | Day | Hour | Keypresses | keypresses_over_last_24_hours |
|-------|------------|------|------------|-------------------------------|
| 30521 | 2024-03-15 | 17 | 3064 | 52104.0 |
| 30522 | 2024-03-15 | 18 | 46 | 45687.0 |
| 30525 | 2024-03-15 | 21 | 13 | 45685.0 |
| 30527 | 2024-03-15 | 23 | 2 | 36963.0 |
| 30538 | 2024-03-16 | 10 | 10299 | 30907.0 |
| 30539 | 2024-03-16 | 11 | 5759 | 34227.0 |
| 30541 | 2024-03-16 | 13 | 862 | 30418.0 |
| 30542 | 2024-03-16 | 14 | 6321 | 34399.0 |
| 30543 | 2024-03-16 | 15 | 4146 | 34674.0 |
| 30544 | 2024-03-16 | 16 | 494 | 31006.0 |
| 30545 | 2024-03-16 | 17 | 1984 | 29926.0 |

Keypresses for the last 48 hours (including hours with 0 keypresses, now that they have been added to our table):

```
[ ]: df_hourly_keypresses.iloc[-48:]
```

```
[ ]:
```

| | Day | Hour | Keypresses | keypresses_over_last_24_hours |
|-------|------------|------|------------|-------------------------------|
| 30498 | 2024-03-14 | 18 | 6463 | 74825.0 |
| 30499 | 2024-03-14 | 19 | 0 | 74825.0 |
| 30500 | 2024-03-14 | 20 | 0 | 72918.0 |
| 30501 | 2024-03-14 | 21 | 15 | 66384.0 |
| 30502 | 2024-03-14 | 22 | 3610 | 66097.0 |
| 30503 | 2024-03-14 | 23 | 5114 | 67124.0 |
| 30504 | 2024-03-15 | 0 | 249 | 65632.0 |
| 30505 | 2024-03-15 | 1 | 0 | 64526.0 |
| 30506 | 2024-03-15 | 2 | 0 | 64526.0 |
| 30507 | 2024-03-15 | 3 | 0 | 64526.0 |
| 30508 | 2024-03-15 | 4 | 0 | 64526.0 |
| 30509 | 2024-03-15 | 5 | 0 | 64526.0 |
| 30510 | 2024-03-15 | 6 | 0 | 64526.0 |
| 30511 | 2024-03-15 | 7 | 0 | 64526.0 |
| 30512 | 2024-03-15 | 8 | 0 | 53658.0 |
| 30513 | 2024-03-15 | 9 | 12936 | 65303.0 |
| 30514 | 2024-03-15 | 10 | 3170 | 64206.0 |

| | | | | |
|-------|------------|----|-------|---------|
| 30515 | 2024-03-15 | 11 | 2439 | 61167.0 |
| 30516 | 2024-03-15 | 12 | 3155 | 58484.0 |
| 30517 | 2024-03-15 | 13 | 1516 | 59214.0 |
| 30518 | 2024-03-15 | 14 | 2340 | 54766.0 |
| 30519 | 2024-03-15 | 15 | 3871 | 53178.0 |
| 30520 | 2024-03-15 | 16 | 4162 | 50036.0 |
| 30521 | 2024-03-15 | 17 | 3064 | 52104.0 |
| 30522 | 2024-03-15 | 18 | 46 | 45687.0 |
| 30523 | 2024-03-15 | 19 | 0 | 45687.0 |
| 30524 | 2024-03-15 | 20 | 0 | 45687.0 |
| 30525 | 2024-03-15 | 21 | 13 | 45685.0 |
| 30526 | 2024-03-15 | 22 | 0 | 42075.0 |
| 30527 | 2024-03-15 | 23 | 2 | 36963.0 |
| 30528 | 2024-03-16 | 0 | 0 | 36714.0 |
| 30529 | 2024-03-16 | 1 | 0 | 36714.0 |
| 30530 | 2024-03-16 | 2 | 0 | 36714.0 |
| 30531 | 2024-03-16 | 3 | 0 | 36714.0 |
| 30532 | 2024-03-16 | 4 | 0 | 36714.0 |
| 30533 | 2024-03-16 | 5 | 0 | 36714.0 |
| 30534 | 2024-03-16 | 6 | 0 | 36714.0 |
| 30535 | 2024-03-16 | 7 | 0 | 36714.0 |
| 30536 | 2024-03-16 | 8 | 0 | 36714.0 |
| 30537 | 2024-03-16 | 9 | 0 | 23778.0 |
| 30538 | 2024-03-16 | 10 | 10299 | 30907.0 |
| 30539 | 2024-03-16 | 11 | 5759 | 34227.0 |
| 30540 | 2024-03-16 | 12 | 0 | 31072.0 |
| 30541 | 2024-03-16 | 13 | 862 | 30418.0 |
| 30542 | 2024-03-16 | 14 | 6321 | 34399.0 |
| 30543 | 2024-03-16 | 15 | 4146 | 34674.0 |
| 30544 | 2024-03-16 | 16 | 494 | 31006.0 |
| 30545 | 2024-03-16 | 17 | 1984 | 29926.0 |

Making sure that all rows with the same day and hour (e.g. from multiple WhatPulse databases) have been merged into the same row:

```
[ ]: df_hourly_keypresses[df_hourly_keypresses.duplicated(
    subset = ['Day', 'Hour'], keep = False)]
```

```
[ ]: Empty DataFrame
Columns: [Day, Hour, Keypresses, keypresses_over_last_24_hours]
Index: []
```

Most keypresses typed in a single hour within the entire dataset:

```
[ ]: df_top_hourly_keypresses = df_hourly_keypresses.sort_values('Keypresses',
    ascending = False).head(50).copy()
df_top_hourly_keypresses['Day and Hour'] = (
    df_top_hourly_keypresses['Day'].astype('str') + ' ' +
```

```
df_top_hourly_keypresses['Hour'].astype('str'))
df_top_hourly_keypresses
```

```
[ ]:
      ID      Day  Hour  Keypresses  keypresses_over_last_24_hours  \
28799  2024-01-03    23      43726                94380.0
28390  2023-12-17    22      38627                39714.0
28558  2023-12-24    22      38424                43794.0
27685  2023-11-18    13      37609                62428.0
27543  2023-11-12    15      34112                45511.0
30274  2024-03-05    10      32901                44898.0
27211  2023-10-29    19      31340                42575.0
27713  2023-11-19    17      30905                79732.0
28584  2023-12-26     0      29522                42802.0
27378  2023-11-05    18      28089                28089.0
28482  2023-12-21    18      27957                55127.0
28596  2023-12-26    12      25519                74109.0
29196  2024-01-20    12      24115                49887.0
28764  2024-01-02    12      22794                48326.0
27864  2023-11-26     0      22455                46258.0
28199  2023-12-09    23      21985                23038.0
29856  2024-02-17     0      21746                34242.0
28914  2024-01-08    18      21473                97706.0
28029  2023-12-02    21      21077                75054.0
28271  2023-12-12    23      19645                41109.0
28216  2023-12-10    16      19210                55219.0
28765  2024-01-02    13      19130                67073.0
28918  2024-01-08    22      18512               105967.0
19342  2022-12-05    22      17552                83422.0
28913  2024-01-08    17      17262                76233.0
28137  2023-12-07     9      17013                73449.0
28433  2023-12-19    17      16879                63734.0
30419  2024-03-11    11      16809                16920.0
29760  2024-02-13     0      16774                23483.0
27919  2023-11-28     7      16542                34636.0
27695  2023-11-18    23      16448               104525.0
28919  2024-01-08    23      15866               121833.0
27993  2023-12-01     9      15575                50349.0
8650   2021-09-16    10      15446                64271.0
28026  2023-12-02    18      15286                56203.0
1958   2020-12-11    14      14618                49483.0
28200  2023-12-10     0      14024                37062.0
27767  2023-11-21    23      13719                63528.0
29728  2024-02-11    16      13641                15438.0
13290  2022-03-28    18      13406                44448.0
28679  2023-12-29    23      13351                28345.0
9442   2021-10-19    10      13345                63492.0
28583  2023-12-25    23      13280                13314.0
```

| | | | | |
|-------|------------|----|-------|---------|
| 28936 | 2024-01-09 | 16 | 13043 | 94248.0 |
| 28784 | 2024-01-03 | 8 | 13001 | 88729.0 |
| 28127 | 2023-12-06 | 23 | 12960 | 67447.0 |
| 7907 | 2021-08-16 | 11 | 12960 | 24082.0 |
| 30513 | 2024-03-15 | 9 | 12936 | 65303.0 |
| 28072 | 2023-12-04 | 16 | 12935 | 43742.0 |
| 3378 | 2021-02-08 | 18 | 12777 | 53030.0 |

| | Day | Hour |
|-------|------------|------|
| 28799 | 2024-01-03 | 23 |
| 28390 | 2023-12-17 | 22 |
| 28558 | 2023-12-24 | 22 |
| 27685 | 2023-11-18 | 13 |
| 27543 | 2023-11-12 | 15 |
| 30274 | 2024-03-05 | 10 |
| 27211 | 2023-10-29 | 19 |
| 27713 | 2023-11-19 | 17 |
| 28584 | 2023-12-26 | 0 |
| 27378 | 2023-11-05 | 18 |
| 28482 | 2023-12-21 | 18 |
| 28596 | 2023-12-26 | 12 |
| 29196 | 2024-01-20 | 12 |
| 28764 | 2024-01-02 | 12 |
| 27864 | 2023-11-26 | 0 |
| 28199 | 2023-12-09 | 23 |
| 29856 | 2024-02-17 | 0 |
| 28914 | 2024-01-08 | 18 |
| 28029 | 2023-12-02 | 21 |
| 28271 | 2023-12-12 | 23 |
| 28216 | 2023-12-10 | 16 |
| 28765 | 2024-01-02 | 13 |
| 28918 | 2024-01-08 | 22 |
| 19342 | 2022-12-05 | 22 |
| 28913 | 2024-01-08 | 17 |
| 28137 | 2023-12-07 | 9 |
| 28433 | 2023-12-19 | 17 |
| 30419 | 2024-03-11 | 11 |
| 29760 | 2024-02-13 | 0 |
| 27919 | 2023-11-28 | 7 |
| 27695 | 2023-11-18 | 23 |
| 28919 | 2024-01-08 | 23 |
| 27993 | 2023-12-01 | 9 |
| 8650 | 2021-09-16 | 10 |
| 28026 | 2023-12-02 | 18 |
| 1958 | 2020-12-11 | 14 |
| 28200 | 2023-12-10 | 0 |
| 27767 | 2023-11-21 | 23 |

```

29728  2024-02-11 16
13290  2022-03-28 18
28679  2023-12-29 23
9442   2021-10-19 10
28583  2023-12-25 23
28936  2024-01-09 16
28784  2024-01-03 8
28127  2023-12-06 23
7907   2021-08-16 11
30513  2024-03-15 9
28072  2023-12-04 16
3378   2021-02-08 18

```

```
[ ]: df_top_hourly_keypresses.head(3)
```

```
[ ]:
      Day  Hour  Keypresses  keypresses_over_last_24_hours \
28799  2024-01-03    23      43726                    94380.0
28390  2023-12-17    22      38627                    39714.0
28558  2023-12-24    22      38424                    43794.0

      Day and Hour
28799  2024-01-03 23
28390  2023-12-17 22
28558  2023-12-24 22

```

```
[ ]: fig_top_hourly_keypresses = px.bar(df_top_hourly_keypresses, x = 'Day and Hour',
y = 'Keypresses', text_auto = '.0f',
title = 'Top Hourly Keypresses')
fig_top_hourly_keypresses.update_xaxes(type='category')
save_chart(fig_top_hourly_keypresses, 'top_hourly_keypresses')
fig_top_hourly_keypresses

```

Average keypresses by hour:

```
[ ]: df_hourly_pivot = df_hourly_keypresses.pivot_table(index = 'Hour',
values = 'Keypresses', aggfunc = ['mean', 'sum']).reset_index()
df_hourly_pivot.columns = df_hourly_pivot.columns.to_flat_index()
# At this point, the columns will read: "(hour, ), (mean, Keypresses),
↪ and
# (sum, Keypresses)". We could convert them to regular columns via a loop,
# but since there are only 3, the simplest solution is to simply rename them
# as follows:
df_hourly_pivot.columns = ['Hour', 'Average Keypresses', 'Keypresses']

# Determining the percentage of total keypresses typed each hour:
total_keypresses = df_hourly_pivot['Keypresses'].sum()
df_hourly_pivot['pct_of_total'] = 100* df_hourly_pivot[

```



```

    'Keypresses'] / total_keypresses
df_hourly_pivot

df_hourly_pivot

```

```

[ ]:
    Hour  Average Keypresses  Keypresses  pct_of_total
0      0          1039.024352    1322678      4.097857
1      1           484.340927     616566      1.910215
2      2          186.043991     236834      0.733748
3      3           62.025923      78959      0.244627
4      4          16.362137      20829      0.064531
5      5           1.282011       1632      0.005056
6      6           0.265515        338      0.001047
7      7          15.991359      20357      0.063069
8      8          187.118617     238202      0.737986
9      9          762.574234     970757      3.007552
10     10         1401.190888     1783716     5.526222
11     11         1688.091123     2148940     6.657742
12     12         1637.648861     2084727     6.458800
13     13         1826.571092     2325225     7.203899
14     14         2095.917518     2668103     8.266187
15     15         2147.017282     2733153     8.467722
16     16         2191.411626     2789667     8.642811
17     17         1763.339356     2244731     6.954517
18     18         1423.216195     1810331     5.608680
19     19         1070.830189     1362096     4.219980
20     20          980.749214     1247513     3.864984
21     21         1140.796384     1451093     4.495706
22     22         1561.058176     1985666     6.151894
23     23         1678.615566     2135199     6.615170

```

```

[ ]: sum(df_hourly_pivot['pct_of_total'])
     # Making sure the percentages were calculated correctly
     # (they should add up to 100%)

```

```

[ ]: 100.0

```

3.2 Plotting average keypresses by hour:

Not surprisingly, my hourly keypress averages are highest during the workday and lowest in the middle of the night, although I've been pretty active in the late evening hours also. (More on this below.)

```

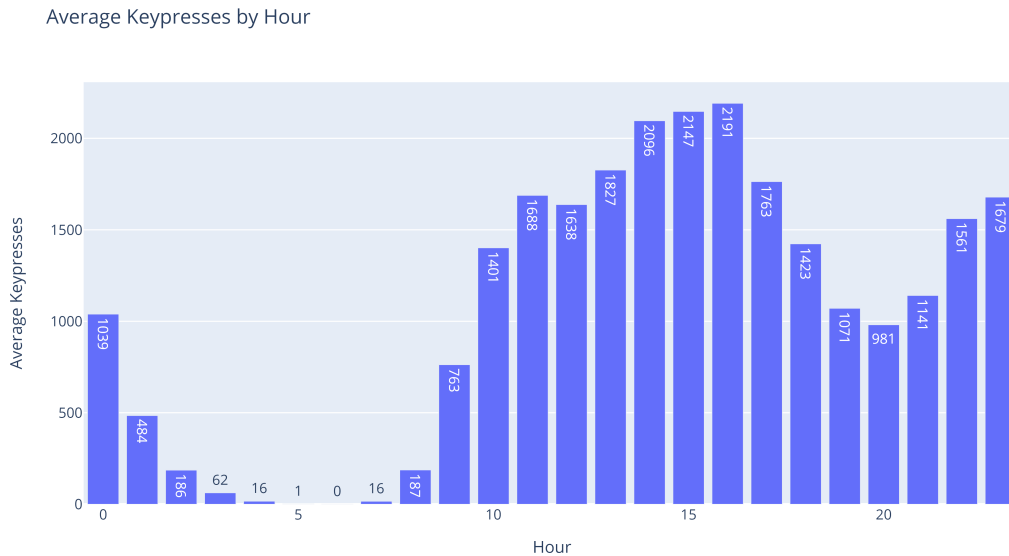
[ ]: fig_hourly_keypresses = px.bar(df_hourly_pivot, x = 'Hour',
    y = 'Average Keypresses', text_auto = '.0f',
    title = 'Average Keypresses by Hour')

```

```
save_chart(fig_hourly_keypresses, 'average_keypresses_by_hour')
fig_hourly_keypresses
```

```
[ ]: Image(static_graphs_folder+'average_keypresses_by_hour.png')
```

```
[ ]:
```



```
[ ]: df_hourly_keypresses['Day'][0]
```

```
[ ]: datetime.date(2020, 9, 21)
```

Saving the updated version of df_hourly_keypresses to a .csv file:

```
[ ]: df_hourly_keypresses.to_csv('data/df_combined_hourly_keypresses_updated.csv')
```

3.3 Data analysis question: Did marriage change my typing patterns?

I got married in April 2023, and I suspected that my keypress distributions as a married man might skew earlier than they did when I was a bachelor. I decided to investigate this graphically by creating subsets of df_hourly_keypresses that contained pre-marriage and post-marriage datasets, then comparing them via a grouped bar chart.

Setting datetime.date() values that will be used to filter df_hourly_keypresses:

```
[ ]: post_mba_work_start_date = datetime.date(2022, 6, 21) # I began my current
# full-time work in June 2022 after finishing my MBA. I chose to limit the
# dataset to this date range so that my results wouldn't be influenced
# by my time as an MBA student (which featured more irregular computer hours).
last_day_before_marriage = datetime.date(2023, 4, 14)
```

```
day_after_honeymoon = datetime.date(2023, 4, 29) # I didn't type much at all
# on my honeymoon, so I excluded this period from my analysis in order not
# to skew the average keypress totals downward.
```

```
[ ]: df_hourly_keypresses_pre_marriage = df_hourly_keypresses.query(
    "Day > @post_mba_work_start_date & Day <= @last_day_before_marriage"
).pivot_table(index = 'Hour', values = 'Keypresses', aggfunc = 'mean').
    ↪reset_index()
df_hourly_keypresses_pre_marriage['Period'] = 'Before Marriage'

df_hourly_keypresses_post_marriage = df_hourly_keypresses.query(
    "Day >= @day_after_honeymoon").pivot_table(
    index = 'Hour', values = 'Keypresses', aggfunc = 'mean').reset_index()
df_hourly_keypresses_post_marriage['Period'] = 'After Marriage'

# Combining these two DataFrames together:
df_hourly_keypresses_by_period = pd.concat([df_hourly_keypresses_pre_marriage,
df_hourly_keypresses_post_marriage])

df_hourly_keypresses_by_period
```

```
[ ]:
  Hour  Keypresses  Period
0     0    641.848485  Before Marriage
1     1    296.060606  Before Marriage
2     2     58.700337  Before Marriage
3     3    31.047138  Before Marriage
4     4     0.090909  Before Marriage
5     5     0.000000  Before Marriage
6     6     1.138047  Before Marriage
7     7     4.380471  Before Marriage
8     8    62.777778  Before Marriage
9     9   574.838384  Before Marriage
10    10  1439.505051  Before Marriage
11    11  1632.949495  Before Marriage
12    12  1639.750842  Before Marriage
13    13  1856.245791  Before Marriage
14    14  1962.851852  Before Marriage
15    15  2019.952862  Before Marriage
16    16  1950.003367  Before Marriage
17    17  1412.973064  Before Marriage
18    18   941.222222  Before Marriage
19    19   888.138047  Before Marriage
20    20   873.659933  Before Marriage
21    21  1083.676768  Before Marriage
22    22  1323.804714  Before Marriage
23    23  1364.552189  Before Marriage
0     0  1038.467492  After Marriage
```

| | | | |
|----|----|-------------|----------------|
| 1 | 1 | 266.263158 | After Marriage |
| 2 | 2 | 112.269350 | After Marriage |
| 3 | 3 | 28.882353 | After Marriage |
| 4 | 4 | 22.250774 | After Marriage |
| 5 | 5 | 0.000000 | After Marriage |
| 6 | 6 | 0.000000 | After Marriage |
| 7 | 7 | 54.566563 | After Marriage |
| 8 | 8 | 271.743034 | After Marriage |
| 9 | 9 | 1212.600619 | After Marriage |
| 10 | 10 | 1878.086687 | After Marriage |
| 11 | 11 | 2082.340557 | After Marriage |
| 12 | 12 | 1938.761610 | After Marriage |
| 13 | 13 | 1939.702786 | After Marriage |
| 14 | 14 | 2022.464396 | After Marriage |
| 15 | 15 | 2046.687307 | After Marriage |
| 16 | 16 | 2070.566563 | After Marriage |
| 17 | 17 | 1874.136223 | After Marriage |
| 18 | 18 | 1551.385093 | After Marriage |
| 19 | 19 | 1125.242236 | After Marriage |
| 20 | 20 | 786.105590 | After Marriage |
| 21 | 21 | 1055.357143 | After Marriage |
| 22 | 22 | 1489.568323 | After Marriage |
| 23 | 23 | 1749.664596 | After Marriage |

My daily keypress counts have increased slightly since getting married (at least when the honeymoon isn't taken into account):

```
[ ]: df_hourly_keypresses_by_period.pivot_table(
      index = 'Period', values = 'Keypresses', aggfunc = 'sum')
```

```
[ ]:
      Keypresses
Period
After Marriage    26617.112455
Before Marriage   22060.168350
```

However, as the following chart shows, the hourly distribution of these keypresses has changed significantly. I'm now typing much less late at night and am getting more keypresses in earlier in the day.

```
[ ]: fig_keypresses_by_period = px.bar(df_hourly_keypresses_by_period, x = 'Hour',
y = 'Keypresses', color = 'Period', barmode = 'group', text_auto = '.0f',
title = 'Average Keypresses by Hour Before and After Getting Married')
save_chart(fig_keypresses_by_period, 'keypresses_before_and_after_marriage')
# See https://plotly.com/python/bar-charts/
# for the use of the 'color' and 'barmode' arguments.
fig_keypresses_by_period
```


c:\Users\kburc\miniforge3\envs\ga15pyd\lib\site-

packages\plotly\express\core.py:1992: FutureWarning:

When grouping with a length-1 list-like, you will need to pass a length-1 tuple to get_group in a future version of pandas. Pass `(name,)` instead of `name` to silence this warning.

```
[ ]: Image(static_graphs_folder+'keypresses_before_and_after_marriage.png')
```

```
[ ]:
```



```
[ ]: end_time = time.time()
run_time = end_time - start_time
run_minutes = run_time // 60
run_seconds = run_time % 60

run_minutes, run_seconds
```

```
[ ]: (0.0, 9.556999921798706)
```

The input() function within the following cell keeps the console window open when running the file in a command prompt. It's not necessary for the Jupyter Notebook, but when I export this notebook as a Python script and then run the script on a scheduled basis, this line gives me time to read the output. See nosklo's response at: <https://stackoverflow.com/a/1000968/13097194>

```
[ ]: print("The program has finished running. Press Enter to exit.") # Lets me know
# that I can now close the program after it has finished running in a console
# window. (I wouldn't want to close it while the
# graphs are in the process of being generated.)
```

```
input()
```

The program has finished running. Press Enter to exit.

```
[ ]: ''
```

That's it for this program! I hope you enjoy using it to analyze your own WhatPulse keypress statistics.